

Introduction

This document aims to compile most (if not all) of the essential Databricks, Apache Spark™, and Delta Lake best practices and optimization techniques in one place. All data engineers and data architects can use it as a guide when designing and developing optimized and cost-effective and efficient data pipelines. It should not be treated as an afterthought but as one of the most important nonfunctional requirements right at the inception of the project. As a result, all of the practices discussed in this book should be considered throughout the development and productionization of data pipelines. This book is divided into several sections, each focusing on a particular problem statement and deep diving into providing all the possible solutions for it.

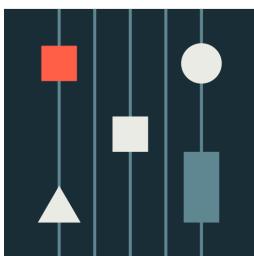


Delta Lake — The Lakehouse Format

Delta Lake is an open format storage layer that delivers reliability, security and performance on your data lake. It enables building a [Lakehouse architecture](#) with compute engines including Spark, PrestoDB, Flink, Trino, and Hive, and APIs for Python, SQL, Scala, Java, Rust, and Ruby. It has many benefits over other formats like Parquet, Avro, ORC, etc. — for example:

- Delta is a protocol to guarantee ACID transactions, high performance and a ton of other features on top of the open source Parquet format. Delta is [open source](#) and the complete protocol can be found [here](#).
- Delta Lake supports ACID transactions and unifies batch and streaming paradigms, simplifying delete/insert/update transactions on incremental data
- Delta allows you to time travel and to read a point-in-time snapshot version of the table
- Delta also comes with many performance enhancements around efficient data layout, indexation, data skipping and caching, etc.

Therefore, it is highly recommended to use Delta as the default data lake storage format to reap all the benefits. On Databricks DBR 8.x and above, Delta is the default format.



Underlying Data Layout

Underneath the hood of a Delta table are Parquet files that store the data. It also contains a subdirectory named `_delta_Log` that stores the Delta transaction log right next to the Parquet files. The size of these Parquet files is really crucial for query performance.

The tiny files problem is a well-known problem in the big data world. When a table has too many underlying tiny files, read latency suffers as a result of the time spent for just opening and closing those tiny files.

To avoid this problem, it's best to have a file size between 16MB and 1GB, which is configurable on a case-by-case basis based on the workload and the specific use case.

If you are using Databricks Runtime 11.3 and above to create managed Delta tables cataloged in [Unity Catalog](#) (Databricks' data catalog), you don't need to worry about optimizing the underlying file sizes or configuring a target file size for your Delta tables because Databricks will carry out this task automatically in the background as part of the auto-tuning capability. In the future, this will also be available for external tables.

For the remaining cases, you would need to manually compact files together using specific Spark jobs to obtain the appropriate file size, but this isn't necessary for Delta tables because they come with out-of-the-box bin packing (compaction) features. In Delta, bin packing can be accomplished in two ways, as detailed below:

1. Optimize & Z-order

`OPTIMIZE` compacts the files to get a file size of up to 1GB, which is [configurable](#). This command basically attempts to size the files to the size that you have configured (or 1GB by default if not configured). You can also combine the `OPTIMIZE` command with the `ZORDER`, which physically sorts or co-locates data across chosen column(s).

SQL (AUTO-DETECTED)

```
OPTIMIZE table_name [WHERE predicate]  
[ZORDER BY (col_name1 [, ...] ) ]
```

- Always choose high cardinality columns (for example: `customer_id` in an `orders` table) for Z-ordering. Date columns are usually low cardinality columns, so they should not be used for Z-order — they are a better fit as the partitioning columns (but you don't always have to partition the tables; please refer to the [Partitioning](#) section below for more details). For Z-order, choose the columns that are most frequently used in filter clauses or as join keys in the downstream queries.
- Never use more than 4 columns since too many columns will degrade Z-ordering effectiveness
- Always run **OPTIMIZE** command on a separate job cluster and not as part of the job itself; otherwise, it might impact the corresponding job's SLA
- Compute-optimized instance family is recommended for the **OPTIMIZE** command as it's a compute-intensive operation
- OPTIMIZE** (with or without **ZORDER**) should be done on a regular basis, such as once a day (or weekly or as per your requirements) to maintain a good file layout for better downstream query performance

2. Auto optimize

[Auto optimize](#), as the name suggests, automatically compacts small files during individual writes to a Delta table, and by default, it tries to achieve a file size of 128MB. It comes with two features:

1. Optimize Write

Optimize Write dynamically optimizes Apache Spark partition sizes based on the actual data, and attempts to write out 128MB files for each table partition. It runs inside the same Spark job.

2. Auto Compact

Following the completion of the Spark job, Auto Compact launches a new job to see if it can further compress files to attain a 128MB file size.

PYTHON (AUTO-DETECTED)

```
-- Table properties  
delta.autoOptimize.optimizeWrite = true  
delta.autoOptimize.autoCompact = true  
  
-- In Spark session conf for all new tables  
set spark.databricks.delta.properties.defaults.autoOptimize.optimizeWrite = true  
set spark.databricks.delta.properties.defaults.autoOptimize.autoCompact = true
```

- Always enable `optimizeWrite` table property if you are not already leveraging the manual **OPTIMIZE** (with or without **ZORDER**) command to get a decent file size. Keep in mind that optimized writes require the shuffling of data according to the partitioning structure of the table. This shuffle naturally incurs additional cost. However, the throughput gains during the write may pay off the cost of the shuffle. The throughput gains when querying the data should still make this feature worthwhile.
- If your job isn't strictly SLA bound, then in that case you can also enable `autoCompact` table property to further take advantage of bin packing

3. Partitioning

[Partitioning](#) can speed up your queries if you provide the partition column(s) as filters or join on partition column(s) or aggregate on partition column(s) or merge partition column(s), as it will help Spark to skip a lot of unnecessary data partitions (i.e., subfolders) during scan time.

- Databricks recommends not to partition tables under 1TB in size and let [ingestion time clustering](#) automatically take effect. This feature will cluster the data based on the order the data was ingested by default for all tables.
- You can partition by a column if you expect data in each partition to be at least 1GB
- Always choose a low cardinality column — for example, year, date — as a partition column
- You can also take advantage of Delta's [generated columns](#) feature while choosing the partition column. Generated columns are a special type of column whose values are automatically generated based on a user-specified function over other columns in the table.

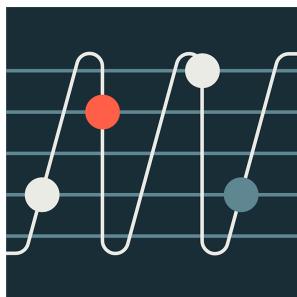
4. File size tuning

In cases where the default file size targeted by Auto-optimize (128MB) or Optimize (1GB) isn't working for you, you can fine-tune it as per your requirement. Set the [target file size](#) by using `delta.targetFileSize` table property and then Auto-optimize and Optimize will binpack to achieve the specified size instead.

```
-- Table properties
delta.targetFileSize = 134217728
```

You can also delegate this task to Databricks if you don't want to manually set the target file size by turning on the `delta.tuneFileSizesForRewrites` table property. When this property is set to true, Databricks will automatically tune the file sizes based on workloads. For example, if you do a lot of merges on the table, then the files will automatically be tuned to much smaller sizes than 1GB to accelerate the merge operation.

```
-- Table properties
delta.tuneFileSizesForRewrite = true
```



Data Shuffling — Why It Happens and How to Control It

Data shuffle occurs as a result of wide transformations such as joins, aggregations and window operations, among others. It is an expensive process since it moves data over the network between worker nodes. We may use a few optimization approaches to either eliminate or improve the efficiency and speed of shuffles.

1. Broadcast hash join

To entirely avoid data shuffling, broadcast one of the two tables or DataFrames (the smaller one) that are being joined together. The table is broadcast by the driver who copies it to all worker nodes.

When executing joins, Spark automatically broadcasts tables less than 10MB; however, we may adjust this threshold to broadcast even larger tables, as demonstrated below:

PYTHON (AUTO-DETECTED)

```
set spark.sql.autoBroadcastJoinThreshold = <size in bytes>
```

When you know that some of the tables in your query are small tables, you can tell Spark to broadcast them explicitly using [hints](#), which is a recommended configuration:

SQL (AUTO-DETECTED)

```
SELECT /*+ BROADCAST(t) */ * FROM <table-name> t
```

Spark 3.0 and above comes with [AQE \(Adaptive Query Execution\)](#), which can also convert the sort-merge join into broadcast hash join (BHJ) when the run time statistics of any join side is smaller than the adaptive broadcast hash join threshold, which is 30MB by default. You can also increase this threshold by changing the following configuration:

PYTHON (AUTO-DETECTED)

```
set spark.databricks.adaptive.autoBroadcastJoinThreshold = <size in bytes>
```

- Note that broadcast hash join is not supported for a full outer join. For the right outer join, only the left-side table can be broadcast, a other left joins only the right-side table can be broadcast.
- If you're running a driver with a lot of memory (32GB+), you can safely raise the broadcast thresholds to something like **200MB**

PYTHON (AUTO-DETECTED)

```
set spark.sql.autoBroadcastJoinThreshold = 209715200;
set spark.databricks.adaptive.autoBroadcastJoinThreshold = 209715200;
```

- Always explicitly broadcast smaller tables using hints or PySpark broadcast function
- Why do we need to explicitly broadcast smaller tables if AQE can automatically broadcast smaller tables for us? The reason for this AQE optimizes queries while they are being executed.
 - a. Spark needs to shuffle the data on both sides and then only AQE can alter the physical plan based on the statistics of the shuffle and convert to broadcast join
 - b. Therefore, if you explicitly broadcast smaller tables using hints, it skips the shuffle altogether and your job will not need to wait for AQE's intervention to optimize the plan
- Never broadcast a table bigger than 1GB because broadcast happens via the driver and a 1GB+ table will either cause OOM on the driver or make the driver unresponsive due to large GC pauses
- Please take note that the size of a table in disk and memory will never be the same. Delta tables are backed by Parquet files, which have varying levels of compression depending on the data. And Spark might broadcast them based on their size in the disk — however they might actually be really big (even more than 8GB) in memory after the decompression and conversion from column to row format. Spark has a hard limit of 8GB on the table size it can broadcast. As a result, your job may fail with an exception in this circumstance case, the solution is to either disable broadcasting by setting `spark.sql.autoBroadcastJoinThreshold` to -1 and do the explicit broadcast using hints (or the PySpark broadcast function) of the tables that are really small in the disk as well as in memory, or set the `spark.sql.autoBroadcastJoinThreshold` to smaller values like 100MB or 50MB instead of setting the threshold to -1.
- The driver can only collect up to 1GB of data in memory at any given time, and anything more than that will trigger an error in the driver causing the job to fail. However, since we want to broadcast tables larger than 10MB, we risk running into this problem. This problem can be solved by increasing the value of the following driver configuration.
 - Please keep in mind that because this is a driver setting; it cannot be altered once the cluster is launched. Therefore, it should be under the cluster's advanced options as a Spark config. Setting this parameter to 8GB for a driver with >32GB memory seems to be fine in most circumstances. In certain cases where the broadcast hash join is going to broadcast a very large table, setting this to 16GB would also make sense.
 - In [Photon](#), we have the executor-side broadcast. So, you don't have to change the following driver configuration if you use a Data Runtime (DBR) with Photon.

PYTHON (AUTO-DETECTED)

```
spark.driver.maxResultSize 16g
```

2. Shuffle hash join over sort-merge join

In most cases Spark chooses sort-merge join (SMJ) when it can't broadcast tables. Sort-merge joins are the most expensive ones. Shuffle-hash join (SHJ) has been found to be faster in some circumstances (but not all) than sort-merge since it does not require an extra sorting step like SMJ. There is a setting that allows you to advise Spark that you would prefer SHJ over SMJ, and with that Spark will try to use SHJ instead of SMJ wherever possible. Please note that this does not mean that Spark will always choose SHJ over SMJ. We are simply defining your preference for this option.

PYTHON (AUTO-DETECTED)

```
set spark.sql.join.preferSortMergeJoin = false
```

Databricks [Photon](#) engine also replaces sort-merge join with shuffle hash join to boost the query performance.



- Setting the `preferSortMergeJoin` config option to false for each job is not necessary. For the first execution of a concerned you can leave this value to default (which is true).
- If the job in question performs a lot of joins, involving a lot of data shuffling and making it difficult to meet the desired SLA, th can use this option and change the `preferSortMergeJoin` value to false

3. Leverage cost-based optimizer (CBO)

Spark SQL can use a [Cost-based optimizer](#) (CBO) to improve query plans. This is especially useful for queries with multiple joins. The CBO is enabled by de You disable the CBO by changing the following configuration:

PYTHON (AUTO-DETECTED)

```
set spark.sql.cbo.enabled = false
```

For CBO to work, it is critical to collect table and column statistics and keep them up to date. Based on the stats, CBO chooses the most economical join str So you will have to run the following SQL command on the tables to compute stats. The stats will be stored in the Hive metastore.

SQL (AUTO-DETECTED)

```
ANALYZE TABLE table_name COMPUTE STATISTICS FOR COLUMNS col1, col2, ...;
```

Join reorder

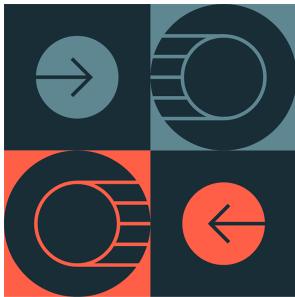
For faster query execution CBO can also use the stats calculated by ANALYZE TABLE command to find the optimal order in which the tables should be joine instance, joining smaller tables first would significantly improve the performance). Join reordering works only for INNER and CROSS joins. To leverage this f set the following configurations:

PYTHON (AUTO-DETECTED)

```
set spark.sql.cbo.enabled = true
set spark.sql.cbo.joinReorder.enabled = true
set spark.sql.statistics.histogram.enabled = true
-- CostBasedJoinReorder requires statistics on the table row count at the very least and its accuracy is improved by having
ANALYZE TABLE table_name COMPUTE STATISTICS FOR COLUMNS col1, col2...;
-- The maximum number of tables in a query for which this joinReorder can be used (default is 12)
set spark.sql.cbo.joinReorder.dp.threshold = <number of tables>
```



- To properly leverage CBO optimizations, ANALYZE TABLE command needs to be executed regularly (preferably once per day or data has mutated by more than 10%, whichever happens first)
- When Delta tables are recreated or overwritten on a daily basis, the ANALYZE TABLE command should be executed immediately the table is overwritten as part of the same job or pipeline. This will have an impact on your pipeline's overall SLA. As a result, in c like this, there is a trade-off between better downstream performance and the current job's execution time. If you don't want CBO optimization to affect the current job's SLA, you can turn it off.
- Never run ANALYZE TABLE command as part of your job. It should be run as a separate job on a separate job cluster. For example, it can be run inside the same nightly notebook running commands like Optimize, Z-order and Vacuum.
- Leverage join reorder where a lot of inner-joins and/or cross-joins are being performed in a single query
- Spark's [Adaptive Query Execution \(AQE\)](#), which changes the query plan on the fly during runtime to a better one, also takes adva the statistics calculated by ANALYZE TABLE command. Therefore, it's recommended to run ANALYZE TABLE command regularly to keep the table statistics updated.



Data Spilling — Why It Happens and How to Get Rid of It

The default setting for the number of Spark SQL shuffle partitions (i.e., the number of CPU cores used to perform wide transformations such as joins, aggregations, etc.) is 200, which isn't always the best value. As a result, each Spark task (or CPU core) is given a large amount of data to process, and if the memory available to each core is insufficient to fit all of that data, some of it is spilled to disk. Spilling to disk is a costly operation, as it involves data serialization, deserialization, reading and writing to disk, etc. Spilling needs to be avoided at all costs and in doing so, we must tune the number of shuffle partitions. There are a couple of ways to tune the number of Spark SQL shuffle partitions as discussed below.

1. AQE auto-tuning

Spark AQE has a feature called `autoOptimizeShuffle` (AOS), which can automatically find the right number of shuffle partitions. Set the following configuration to enable auto-tuning:

PYTHON (AUTO-DETECTED)

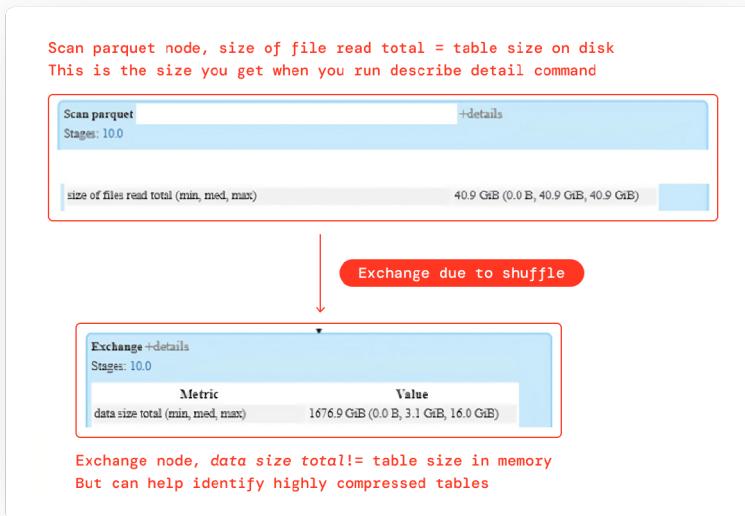
```
set spark.sql.shuffle.partitions=auto
```

Caveat: unusually high compression

There are certain limitations to AOS. AOS may not be able to estimate the correct number of shuffle partitions in some circumstances where source tables have an unusually high compression ratio (20x to 40x).

There are two ways you can identify the highly compressed tables:

1. Spark UI SQL DAG



Although "data size total" metrics in the Exchange node don't provide the exact size of a table in memory, it can definitely help identify the highly compressed table. The Scan Parquet node provides the precise size of a table in the disk. The Exchange node data size in the aforementioned case is 40x larger than the size on the disk, indicating that the table is probably heavily compressed on the disk.

2. Cache the table

A table can be cached in memory to figure out its actual size in memory. Here's how to go about it:

PYTHON (AUTO-DETECTED)

```
-- count here is forcing the cache materialization
spark.table("table").cache().count()
```

Refer to the storage tab of Spark UI to find the size of the table in memory after the command above has been completed:

Storage																
MB	Environment	Executors	SQL	Structured Streaming	JDBC/ODBC Server											
1155.0 MB	300.0 GiB	0 %	310.8 KiB	6.0 GiB	0 %	0.0 B	1155.0 MiB	1155.0 MiB	1104.9 MiB	0.0 B	50.0 MiB	0 %	15			
1118.7 MB	300.0 GiB	0 %	328.9 KiB	6.0 GiB	0 %	0.0 B	1118.7 MiB	1118.7 MiB	1073.5 MiB	0.0 B	45.2 MiB	0 %	15			
1067.3 MB	300.0 GiB	0 %	294.5 KiB	6.0 GiB	0 %	0.0 B	1067.3 MiB	1067.3 MiB	1067.3 MiB	0.0 B	0.0 B	0 %	16			
1128.5 MB	300.0 GiB	0 %	277.1 KiB	6.0 GiB	0 %	0.0 B	1128.5 MiB	1128.5 MiB	1128.5 MiB	0.0 B	0.0 B	0 %	16			
1004.6 MB	300.0 GiB	0 %	276.7 KiB	6.0 GiB	0 %	0.0 B	1004.6 MiB	1004.6 MiB	1004.6 MiB	0.0 B	0.0 B	0 %	16			
17.2 GiB	4.7 TiB	0 %	4.8 MiB	96.0 GiB	0 %	0.0 B	17.2 GiB	17.2 GiB	17.0 GiB	0.0 B	215.2 MiB	0 %	253			
Data Read from IO Cache (Cache Hits, Compressed)			Data Written to IO Cache (Compressed)		Cache Misses (Compressed)	True Cache Misses	Partial Cache Misses	Rescheduling Cache Misses	Cache Hit Ratio	Number of Local Scan Tasks		Number of Rescheduled Scan Tasks				
0.0 B			17.2 GiB		17.2 GiB	17.0 GiB	0.0 B	215.2 MiB	0 %	253		8				
Storage Level						Cached Partitions			Fraction Cached	Size in						
						261			100%	343.9 GiB						

Solution: To counter this effect, reduce the value of the per partition size used by AQE to determine the initial shuffle partition number (default 128MB) as follows:

PYTHON (AUTO-DETECTED)

```
-- setting to 16MB for example
set spark.databricks.adaptive.autoOptimizeShuffle.preshufflePartitionSizeInBytes = 16777216
```

After lowering the preshufflePartitionSizeInBytes value to 16MB, if AOS is still calculating the incorrect number of partitions and you are still experiencing large data spills, you should further lower the preshufflePartitionSizeInBytes value to 8MB. If this still doesn't resolve your spill issue, it is best to disable and manually tune the number of shuffle partitions as explained in the next section.

2. Manually fine-tune

To manually fine-tune the number of shuffle partitions we need:

1. The total amount of shuffled data. To do so, run the Spark query once, and then use the Spark UI to retrieve this value, as demonstrated in the example

Details for Stage 27 (Attempt 0)

Resource Profile Id: 0

Total Time Across All Tasks: 43.9 h

Locality Level Summary: Process local: 1075

Output Size / Records: 235.0 GiB / 4130744580

Shuffle Read Size / Records: 580.0 GiB / 4724958078

← **Total size of shuffle**

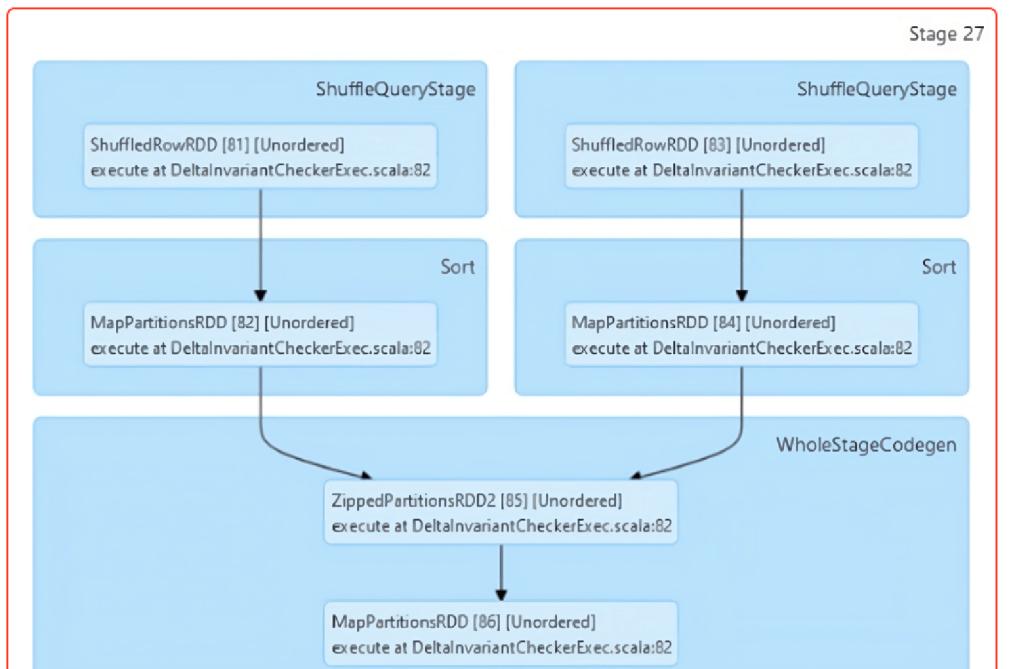
Spill (Memory): 3.0 TiB

Data spill

Spill (Disk): 526.2 GiB

Associated Job Ids: 17

DAG Visualization



- As a rule of thumb, we need to make sure that after tuning the number of shuffle partitions, each task should approximately be processing **128MB to 200** data. You can see this value in the summary metrics for the shuffle stage in Spark UI, as shown in the example below:

Summary Metrics for 30 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0 ms	7 ms	0.1 s	0.3 s	3 s
Scheduler Delay	0 ms	0 ms	1 ms	2 ms	53 ms
Task Deserialization Time	0 ms	1 ms	1 ms	1 ms	7 ms
GC Time	0 ms	0 ms	39 ms	0.2 s	2 s
Peak Execution Memory	0.0 B	79.3 KB	5.5 MB	26.5 MB	274.0 MB
Shuffle Read Blocked Time	0 ms	0 ms	0 ms	0 ms	1 ms
Shuffle Read Size / Records	0.0 B / 0	15.5 KB / 903	1095.3 KB / 69750	4.1 MB / 211991	35.0 MB / 2730493

↑ **Data processed by each task**

So here is the formula to compute the right number of shuffle partitions:

PYTHON (AUTO-DETECTED)

Let's assume that:

Total number of total worker cores **in** cluster = T

Total amount of data being shuffled **in** shuffle stage (**in** megabytes) = B

Optimal size of data to be processed per task (**in** megabytes) = 128

Hence the multiplication factor (M): $M = \text{ceiling}(B / 128 / T)$

And the number of shuffle partitions (N): $N = M \times T$

Note that we have used the ceiling function here to ensure that **all** the cluster cores are fully engaged till the very last

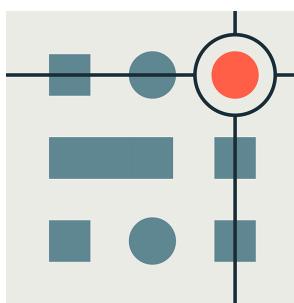
- The optimal size of data to be processed per task should be 128MB approximately. It works well in most cases. It might not work if there is some sort of data explosion happening in your query. You might need to choose a smaller value in that case. We will have a section on data explosion later in this document.
 - If you are neither using auto-tune (AOS) nor manually fine-tuning the shuffle partitions, then as a rule of thumb set this to twice, or thrice the number of total worker CPU cores.



PYTHON (AUTO-DETECTED)

```
-- in SQL
set spark.sql.shuffle.partitions = 2*<number of total worker cores in cluster>
-- in PySpark
spark.conf.set("spark.sql.shuffle.partitions", 2*<number of total worker cores in cluster>)
-- or
spark.conf.set("spark.sql.shuffle.partitions", 2*sc.defaultParallelism)
```

- Because there may be multiple Spark SQL queries in a single notebook, fine-tuning the number of shuffle partitions for each query is a time-consuming task. So, our advice is to fine-tune it for the largest query with the greatest number for the total amount of data shuffled for a shuffle stage and then set that value once for the entire notebook.
 - If there is skewness in data, then fine-tuning the shuffle partitions will not help with data spills. In that case, you should first get rid of data skew. Please refer to the next section on data skew for more details.



Data Skewness — Identification and Remediation

Data skew is the situation where only a few CPU cores wind up processing a huge amount of data due to uneven data distribution. For example, when you join aggregate using the columns(s) around which data is not uniformly distributed, then you will end up with a skewed shuffle stage that will take a lot of time to complete (might actually fail as well after several attempts).

Identification of skew

- If all the Spark tasks for the shuffle stage are finished and just one or two of them are hanging for a long time, that's an indication of skew. You can also get information from Spark UI.



	Submitted	Duration	Tasks: Succeeded/Total	Input	Output
+details	2018/02/18 15:12:11	47 h	199/200		

One stuck task

- In the tasks summary metrics, if you see a huge difference between the min and max shuffle read size, that's also an indication of data skewness

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	29 s	1.2 min	1.5 min	1.7 min	4.0 min
GC Time	77.0 ms	0.8 s	1 s	1 s	3 s
Spill (memory)	2.8 GiB	6.9 GiB	8.5 GiB	10.4 GiB	19.2 GiB
Spill (disk)	388.5 MiB	1.2 GiB	1.6 GiB	2.2 GiB	4.3 GiB
Shuffle Read Size / Records	479.6 MiB / 5349999	749.9 MiB / 6923474	1015 MiB / 8417393	1.3 GiB / 9910309	2.3 GiB / 15914688
Shuffle Write Size / Records	477.2 MiB / 5124404	746.3 MiB / 6697251	1008.6 MiB / 8191021	1.2 GiB / 9683993	2.3 GiB / 15688494
Scheduler Delay	4.0 ms	6.0 ms	21.0 ms	29.0 ms	0.7 s
Task Deserialization Time	4.0 ms	7.0 ms	9.0 ms	11.0 ms	29.0 ms
Shuffle Read Blocked Time	0.0 ms	18 s	32 s	51 s	3.1 min
Shuffle Remote Reads	23 MiB	746.1 MiB	1012.4 MiB	1.3 GiB	2.3 GiB
Result Serialization Time	0.0 ms	0.0 ms	0.0 ms	0.0 ms	0.0 ms
Getting Result Time	0.0 ms	0.0 ms	0.0 ms	0.0 ms	0.0 ms
Peak Execution Memory	6 GiB	6.1 GiB	6.3 GiB	7.1 GiB	16.1 GiB

Tasks summary metrics: huge diff between min and max shuffle size

- Even after fine-tuning the number of shuffle partitions, if there are a lot of data spills, then this might actually be because of skewness
- Lastly, you can just simply count the number of rows while grouping by join or aggregation columns. If there is an enormous disparity between the row counts, then it's a definite skew.

SQL (AUTO-DETECTED)

```
SELECT COLUMN_NAME,
COUNT(*)
FROM TABLE
GROUP BY COLUMN_NAME
```

Skew remediation

1. Filter skewed values

If it's possible to filter out the values around which there is a skew, then that will easily solve the issue. If you join using a column with a lot of null values for example, you'll have data skew. In this scenario, filtering out the null values will resolve the issue.

2. Skew hints

In the case where you are able to identify the table, the column, and preferably also the values that are causing data skew, then you can explicitly tell Spark about it using [skew hints](#) so that Spark can try to resolve it for you.

SQL (AUTO-DETECTED)

```
SELECT /*+ SKEW('table', 'column_name', (value1, value2)) */ * FROM table
```

3. AQE skew optimization

Spark 3.0+'s [AQE](#) can also dynamically solve the data skew for you. It's by default enabled, but if you want to disable it, then set the following configuration to false:

PYTHON (AUTO-DETECTED)

```
set spark.sql.adaptive.skewJoin.enabled = false
```

By default any partition that has at least 256MB of data and is at least 5 times bigger in size than the average partition size will be considered as a skew partition by AQE. You can also change these values to fine-tune default AQE behavior:

PYTHON (AUTO-DETECTED)

```
-- default is 5
set spark.sql.adaptive.skewJoin.skewedPartitionFactor = <value>
-- default is 256MB
set spark.sql.adaptive.skewJoin.skewedPartitionThresholdInBytes = <size in bytes>
```

4. Salting

If none of the above-mentioned options work for you, the only other option is to do salting. It's a strategy for breaking a large skewed partition into smaller partitions by appending random integers as suffixes to skewed column values. Please see this example [notebook](#) (from cmd7 onwards) to fix the data skew problem using the salting technique. This example shows how to fix data skew in a join query.

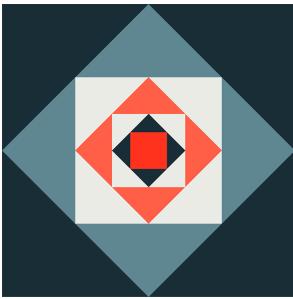
Now let's imagine you have a skewed aggregation query. We can still use salting to remedy the data skew, but in this situation, we would only need to enable partial salting for the value(s) that are contributing to the skewness. Let's assume that for an e-commerce organization, we want to find out the last transaction date for all its customers and the data is skewed due to uneven distribution of the *customer_id* value xyz. So to fix this skewness, we will perform the partial salting as follows:

```
Select
  Case
    When
      salted_key like '%_salt'
    Then
      'xyz'
    Else
      salted_key
  End
  As customer_id, max(transaction_date) as last_transaction_date
From
(
  Select
    Case
      Case
        When
          customer_id is 'xyz'
        Then
          concat(cast( 1 * FLOOR(RAND(12345)*500) AS string), '_', 'salt')
        Else
          customer_id
      End
      As salted_key, max(transaction_date) as transaction_date
    From
      customer_sales_table
    Group by
      salted_key
)
Group by customer_id
```

The diagram illustrates the execution flow of the query. On the right, a large red bracket groups the entire query from 'Select' to 'Group by' under the heading 'Final Aggregation'. On the left, another large red bracket groups the subquery starting with 'From' under the heading 'Intermediate Aggregation'.

Try these solutions in the order they are mentioned above. That means salting should be the last choice, not the first, as it requires schema changes. Hints and AQE solutions are much simpler to implement.

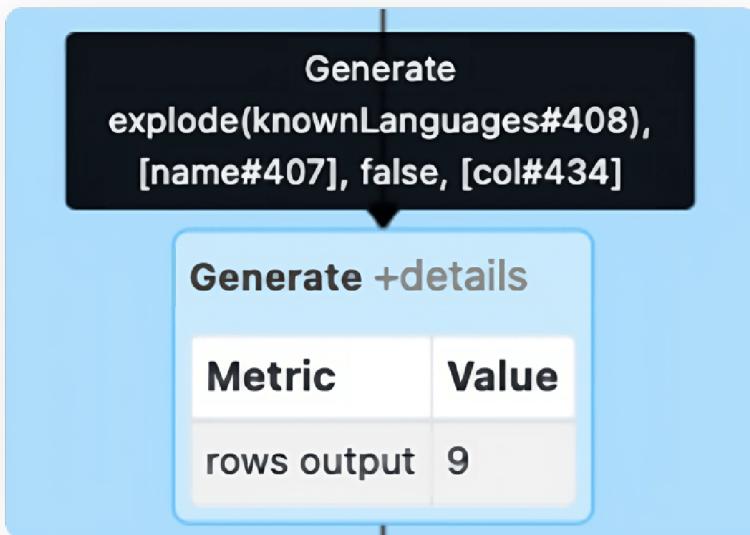




Data Explosion — Identification, Consequences and Solutions

During a Spark job execution after a certain transformation step, you may see an unusually big rise in data volume, which is considered a data explosion. The execution is significantly slowed as a result of this. The following are some of the most prevalent transformations that can result in a data explosion:

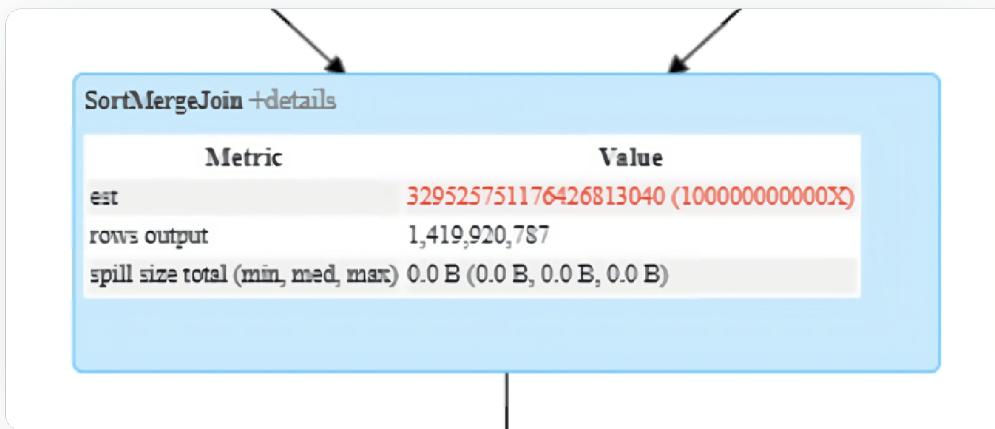
1. Explode function



While working with structured files with the formats like JSON, Parquet, Delta and XML, we often get data in collections like arrays, lists and maps. In such cases, the `explode()` function is useful to convert collection columns to rows in order to process in Spark effectively. This `explode` operation can significantly increase the data volume. The explode operation is represented by the `Generate` node in Spark UI as shown in the image above.

2. Join operation

A common cause of slow queries is joins that produce significantly more rows than anticipated. This is often referred to as a “row explosion.” Refer to the `row output` metric in the `SortMergeJoin` or `ShuffleHashJoin` node of Spark UI to identify a potential row explosion.



Impact

While reading input data from sources like Parquet, Delta, etc., Spark approximately reads 128MB per task per core, which is a very decent partition size to fit the memory available to each CPU core. But due to data explosion, the 128MB data might get converted into a significantly high volume (for example, a couple of GBs), which is problematic as a single CPU core might not have enough memory to fit that exploded partition. As a consequence, for subsequent wide transformations, a lot of data might get spilled into the disk, which can significantly impact the query performance.

Solutions

- Decrease the size of input partitions, i.e., `spark.sql.files.maxPartitionBytes` (default 128MB), to create smaller input partitions in order to counter effect of `explode()` function. Instead of 128MB, you can choose a much smaller partition size like 16MB or 32MB, for example.

PYTHON (AUTO-DETECTED)

```
set spark.sql.files.maxPartitionBytes = <size in bytes>
```

- You can explicitly call the `repartition()` function right after the read statement to increase the total number of partitions. This will allow you to reduce the size of each partition.
- When the explosion is happening due to a join operation, a simple solution would be to increase the number of shuffle partitions, which will decrease the size of the partition to much less than 128MB. Refer to the [manual shuffle partitions tuning](#) section for more information.



Data Skipping and Pruning

The amount of data to process has a direct relationship with query performance. Therefore, it's extremely important to read only the required data and skip a unnecessary data. There are a couple of data skipping and pruning techniques that you can apply with Spark and Delta.

1. Delta data skipping

Delta [data skipping](#) automatically collects the stats (min, max, etc.) for the first 32 columns for each underlying Parquet file when you write data into a Delta table. Databricks takes advantage of this information (minimum and maximum values) at query time to skip unnecessary files in order to speed up the queries.

To collect stats for more than the 32 first columns, you can set the following Delta property:

```
-- table property
delta.dataSkippingNumIndexedCols = <value>
```

Collecting statistics on long strings is an expensive operation. To avoid collecting statistics on long strings, you can either configure the table property `delta.dataSkippingNumIndexedCols` to avoid columns containing long strings or move columns containing long strings to a column greater than `delta.dataSkippingNumIndexedCols` using [ALTER TABLE](#) as shown below:

SQL (AUTO-DETECTED)

```
ALTER TABLE table_name ALTER [COLUMN] col_name col_name data_type [COMMENT col_comment] [FIRST|AFTER colA_name]
```

2. Column pruning

When reading a table, we generally choose all of the columns, but this is inefficient. To avoid scanning extraneous data, always inquire about what columns are part of the workload computation and are needed by downstream queries. Only those columns should be selected from the source database. This has the potential to have a significant impact on query performance.

PYTHON (AUTO-DETECTED)

```
-- SQL
SELECT col1, col2, .. coln FROM table

-- PySpark
dataframe = spark.table("table").select("col1", "col2", ... "coln")
```

3. Predicate pushdown

This aims at pushing down the filtering to the “bare metal” — i.e., a data source engine. That is to increase the performance of queries since the filtering is performed at a very low level rather than dealing with the entire data set after it has been loaded to Spark’s memory.

To leverage the predicate pushdown, all you need to do is add filters when reading data from source tables. Predicate pushdown is data source engine dependent. It works for data sources like Parquet, Delta, Cassandra, JDBC, etc., but it will not work for data sources like text, JSON, XML, etc.

SQL (AUTO-DETECTED)

```
-- SQL
SELECT col1, col2 .. coln FROM table WHERE col1 = <value>

-- PySpark
dataframe = spark.table("table").select("col1", "col2", ... "coln").filter(col("col1") = <value>)
```

In cases where you are performing join operations, apply filters before joins. As a rule of thumb, apply filter(s) right after the table read statement.

4. Partition pruning

The partition elimination technique allows optimizing performance when reading folders from the corresponding file system so that the desired files only in the specified partition can be read. It will address shifting the filtering of data as close to the source as possible to prevent keeping unnecessary data in memory and aim of reducing disk I/O.

To leverage partition pruning, all you have to do is provide a filter on the column(s) being used as table partition(s).

SQL (AUTO-DETECTED)

```
-- SQL
SELECT * FROM table WHERE partition_col = <value>

-- PySpark
dataframe = spark.table("table").filter(col("partition_col") = <value>)
```

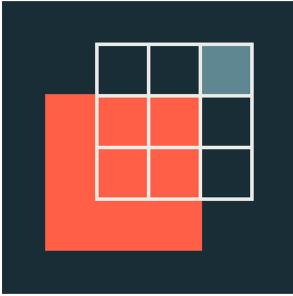
In cases where you are performing join operations, apply partition filter(s) before joins. As a rule of thumb, apply partition filter(s) right after the table read statement.

5. Dynamic partition pruning (DPP)

In Apache Spark 3.0+, a new optimization called Dynamic Partition Pruning (DPP) is implemented. DPP occurs when the optimizer is unable to identify at partition level which partitions it has to eliminate. In particular, we consider a star schema that consists of one or multiple fact tables referencing any number of dimension tables. For such join operations, we can prune the partitions the join reads from a fact table by identifying those partitions that result from filtering the dimension tables. To leverage this feature, no configuration is required. It's enabled by default on Spark 3.0+.

6. Dynamic file pruning (DFP)

Dynamic File Pruning (DFP) is available on Databricks Runtime and is by default enabled on all recent runtimes. As the name suggests, it works in a similar way to DPP, but it performs dynamic pruning at the file level instead of the partition level to further speed up your queries. To leverage this feature, no configuration is required. DFP is automatically enabled in Databricks Runtime 6.1 and higher.



Data Caching — Leverage Caching to Speed Up the Workload

Delta cache and Spark cache are the two different types of caching that you can leverage to make your workloads faster.

1. Delta cache

The [Delta cache](#) accelerates data reads by creating copies of remote files in nodes' local storage (SSD drives) using a fast intermediate data format.

- Use Accelerated instances with Delta cache enabled by default (such as *Standard_E16ds_v4* in the memory-optimized category and *Standard_L4s* in storage-optimized category in Azure cloud)
- Delta cache can be enabled on the workers of other instance families with SSD drives (such as *Fsv2-series* in the compute-optimized category in Azure cloud). To explicitly enable the Delta caching, set the following configuration:

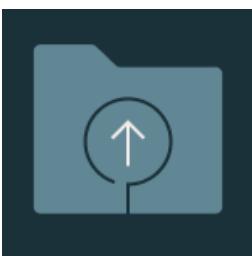
PYTHON (AUTO-DETECTED)

```
set spark.databricks.io.cache.enabled = true
```

2. Spark cache

Using `cache()` and `persist()` methods, Spark provides an optimization mechanism to [cache](#) the intermediate computation of a Spark DataFrame so they can be reused in subsequent actions. Similarly, you can also cache a table using the [CACHE TABLE](#) command. There are different [cache modes](#) that allow you to control where to store the cached data (in the memory, in the disk, in the memory and the disk, with or without serialization, etc.).

- Spark caching is only useful when more than one Spark action (for instance, `count`, `saveAsTable`, `write`, etc.) is being executed on the same DataFrame
-  Databricks recommends using Delta caching instead of Spark caching, as Delta caching provides better performance outcomes. Data stored in the disk cache can be read and operated on faster than the data in the Spark cache. This is because the disk cache uses efficient decompression algorithms and outputs data in the optimal format for further processing using whole-stage code generation.
- Any compute-heavy workload (with a lot of wide transformations like joins and aggregations) would benefit from Delta caching, especially when reading from the same tables over and over. Always use Delta cache-enabled instances for those workloads.



Intermediate Results — When to Persist Them

In large pipelines spanning multiple SQL queries, we often tend to create one or many intermediate working (or temp or staging) Delta tables. This allows us to break a big query into small ones for more readability and maintainability. This strategy, however, has a negative impact on the execution duration of our job.

- We write the staging tables to the cloud storage, which takes time
- Then we read back the staging tables from cloud storage in subsequent steps, which also takes additional time

Therefore, the better approach would be to [create temp views](#) instead of materialized Delta tables since temp views are lazily evaluated and are actually not materialized.

SQL (AUTO-DETECTED)

```
CREATE OR REPLACE TEMP VIEW <view-name> AS SELECT col1, col2, ... FROM <table-name>
```

- Always turn an intermediate table into a temporary view if it is used only once in the same Spark job
- If an intermediate table is being used more than once in the same Spark job, you should leave it as a Delta table rather than turning a temporary view since using a temporary view can cause Spark to execute the related query in part more than once. It is strongly advised to use Delta cache-enabled worker instances in this situation so that the temporary table can be automatically cached in SSD drives to speed up the scanning.



Delta Merge — Let's Speed It Up

You can upsert data from a source table, view or DataFrame into a target Delta table by using the [MERGE](#) operation. Delta Lake supports inserts, updates and deletes in the MERGE command.

Instead of overwriting and inserting the entire Delta table on a daily basis, it is recommended that we use an incremental load strategy wherever possible. To achieve the incremental load, Delta MERGE is very essential. Delta merge can also be used to create SCD Type 2 tables and change data capture (CDC) use cases. Here are a couple of examples of how you can use Delta merge: [SCD Type 2 using Merge](#), [Change data capture using Merge](#).

Merge internals

Behind-the-scenes merge operation is performed in two steps:

1. Inner-join between source and target using the condition in the ON clause to return a list of files in target that contain matching rows to the driver.
2. This step comprises two possibilities — depending on the conditions, one of the following will be executed:
 - a. If there were no matching rows in the target in step 1, an append-only write is performed to append the source to target
 - b. Else a full outer join is performed to consolidate the changes between the source being merged and the matching files produced in step 1.

Merge challenges

There are certain performance issues that can be encountered because of the following reasons:

- If we do not have a specific enough condition to match in the ON clause of the merge operation, we will end up rewriting a lot of data. This can slow down the merge.
- If we rewrite a large amount of data after each merge, the Z-order sorting will be messed up, and we may need to execute Z-order after each merge to restore the sorting.

Merge optimizations

We can use the following optimization techniques to resolve the above-mentioned challenges:

Target table data layout

If the target table contains large files (for example, 500MB-1GB), many of those files will be returned to the drive during step 1 of the merge operation, as the larger the file, the greater the chance of finding at least one matching row. And because of this, a lot of data will be rewritten. Therefore, for the merge-heavy tables recommended to have smaller file sizes from 16MB to 64MB, varying on a case-by-case basis and workload. Refer to the section on [file size tuning](#) for more information.

Partition pruning

Provide the partition filters in the ON clause of the merge operation to discard the irrelevant partitions. Refer to this [example](#) for more details.

File pruning

Provide Z-order columns (if any) as filters in the ON clause of the merge operation to discard the irrelevant files. You can add multiple conditions using the ALL operator in the ON clause.

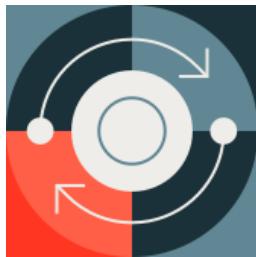
Broadcast join

Since Delta merge performs joins behind the scene, you can speed it up by explicitly broadcasting the source DataFrame to be merged in the target Delta table if the source DataFrame is small enough (<= 200MB). Refer to the [broadcast join](#) section for more information.

Low shuffle merge

- This is a new MERGE algorithm that aims to maintain the existing data organization (including Z-order clustering) for unmodified data, while simultaneously improving performance
- With this “low shuffle” MERGE, only updated rows will be reorganized by the operation, while unchanged rows remain in the same order and file groupings they were in before the operation

Low shuffle merge is enabled by default in Databricks Runtime 10.4 and above. In earlier supported Databricks Runtime versions, it can be enabled by setting configuration `spark.databricks.delta.merge.enableLowShuffle` to true.



Data Purging — What to Do With Stale Data

Delta comes with a [VACUUM](#) feature to purge unused older data files. It removes all files from the table directory that are not managed by Delta, as well as any files that are no longer in the latest state of the transaction log for the table and are older than a retention threshold. The default threshold is 7 days.

```
VACUUM table_name
```

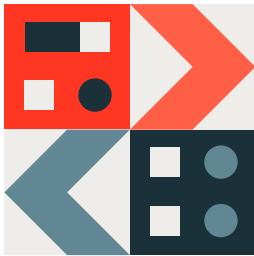
You can change the default retention threshold if needed as follows:

```
-- table properties
deltaTable.deletedFileRetentionDuration = "interval 15 days"
```

VACUUM will skip all directories that begin with an underscore (_), which includes the `_delta_log`. Log files are deleted automatically and asynchronously after checkpoint operations. The default retention period of log files is 30 days, configurable through the `delta.logRetentionDuration`.

- It is recommended that you set a retention interval to be at least 7 days because old snapshots and uncommitted files can still be visible to concurrent readers or writers to the table. If VACUUM cleans up active files, concurrent readers can fail, or, worse, tables can become corrupted when VACUUM deletes files that have not yet been committed.
- Set `deltaTable.deletedFileRetentionDuration` and `delta.logRetentionDuration` to the same value to have the same retention for data and transaction history
- Run the VACUUM command on a daily/weekly basis depending on the frequency of transactions applied on the Delta tables
- Never run this command as part of your job but run it as a separate job on a dedicated job cluster, usually clubbed with the OPTIMIZE command
- VACUUM is not a very intensive operation (the main task is file listing, which is done in parallel on the workers; the actual file deletion is handled by the driver), so a small autoscaling cluster of 1 to 4 workers is sufficient

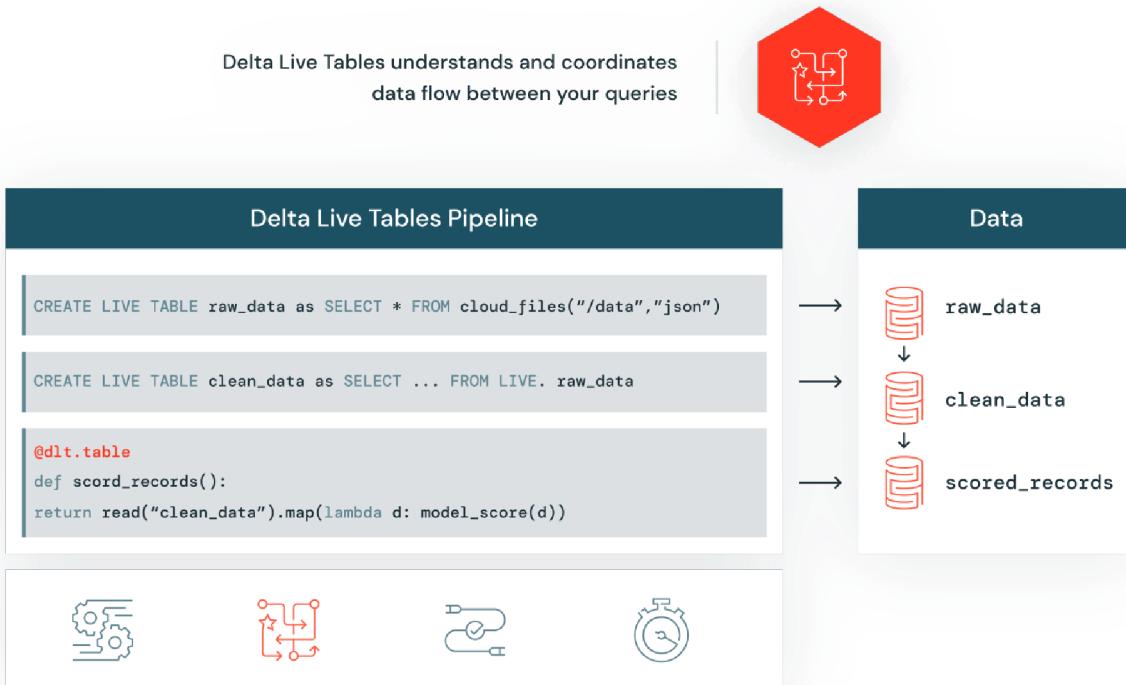




Delta Live Tables (DLT)

Delta Live Tables (DLT) makes it easy to build and manage reliable batch and streaming data pipelines that deliver high-quality data on the [Databricks Lake Platform](#).

DLT helps data engineering teams simplify ETL development and management with declarative pipeline development, automatic data testing, and deep visit monitoring and recovery.



Easily build and maintain data pipelines

With Delta Live Tables, easily define end-to-end data pipelines in SQL or Python. Simply specify the data source, the transformation logic and the destination of the data — instead of manually stitching together siloed data processing jobs. Automatically maintain all data dependencies across the pipeline and reuse pipelines with environment-independent data management. Run in batch or streaming mode and specify incremental or complete computation for each table

Delta Live Tables understands and coordinates data flow between your queries



Delta Live Tables Pipeline

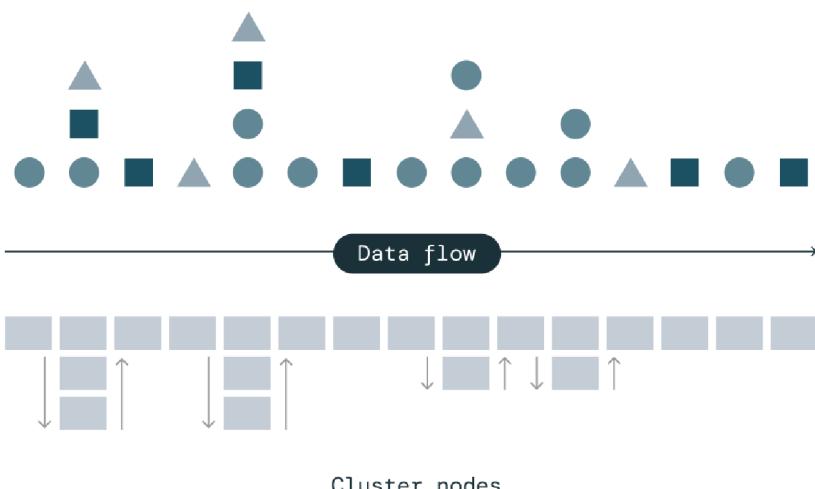
```
CREATE LIVE TABLE raw_data as SELECT * FROM cloud_files("/data","json")  
  
CREATE LIVE TABLE clean_data as SELECT ... FROM LIVE. raw_data  
  
@dlt.table  
def scord_records():  
    return read("clean_data").map(lambda d: model_score(d))
```

Data



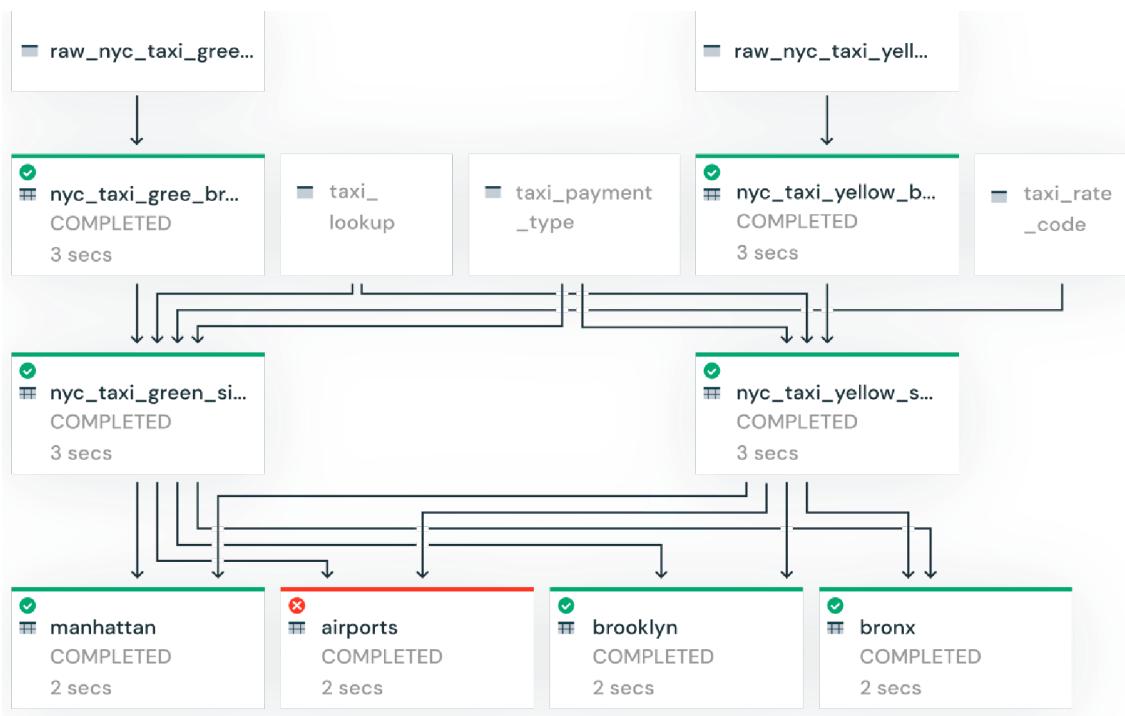
Automatic data quality testing

Delta Live Tables helps to ensure accurate and useful BI, data science and machine learning with high-quality data for downstream users. Prevent bad data flowing into tables through validation and integrity checks and avoid data quality errors with predefined error policies (fail, drop, alert or quarantine data). In a you can monitor data quality trends over time to get insight into how your data is evolving and where changes may be necessary.



Cost-effective streaming through efficient compute autoscaling

Delta Live Tables Enhanced Autoscaling is designed to handle streaming workloads that are spiky and unpredictable. It optimizes cluster utilization by only scaling up to the necessary number of nodes while maintaining end-to-end SLAs and gracefully shuts down nodes when utilization is low to avoid unnecessary spending.

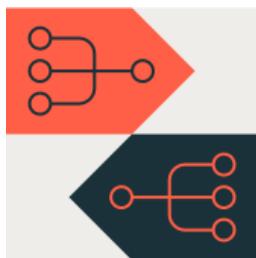


Deep visibility for pipeline monitoring and observability

Gain deep visibility into pipeline operations with tools to visually track operational stats and data lineage. Reduce downtime with automatic error handling and replay. Speed up maintenance with single-click deployment and upgrades.



- DLT is a managed ETL framework that frees users to concentrate on writing business-related code while DLT handles operational tasks such as cluster management, autoscaling, retries on failures, and other configuration tweaking in the background. Therefore, it is advised to use DLT for data engineering workloads as much as possible so that users are spared from worrying about Spark tuning, cluster configuration, and maintenance (see the following part), Delta maintenance tasks, etc.



Databricks Cluster Configuration and Tuning

All-purpose clusters vs. automated clusters

- All-purpose clusters should only be used for ad hoc query execution and interactive notebook execution during the development and/or testing phases
- Never use an all-purpose cluster for an automated job; instead, use ephemeral (also called automated) job clusters for jobs
- This is the single biggest cost optimization impact: you could reduce the [DBU cost](#) by opting for an automated cluster instead of an all-purpose cluster
- You can also leverage [SQL warehouses](#) to run your SQL queries. A SQL warehouse is a compute resource that lets you run [SQL commands](#) on data objects within [Databricks SQL](#). SQL warehouses are also available in [serverless](#) flavor, offering you access to an instant compute.

Autoscaling

Databricks provides a unique feature of [cluster autoscaling](#). Here are some guidelines on when and how to leverage autoscaling:

- Always use autoscaling when running ad hoc queries, interactive notebook execution, and developing/testing pipelines using interactive clusters with minimum workers set to 1. The combination of both can bring good cost savings.
- For an autoscaling job cluster in production, don't set the minimum instances to 1 if the job certainly needs more resources than 1 worker. Instead, set it to a bigger value based on the minimum compute requirements. It will save you some scale-up time.

- You don't necessarily need to employ autoscaling if you've fine-tuned the Spark shuffle partitions to use all of the worker cores for a given job and that justify its own job cluster. This is the best option because it allows complete cost control. However, you can still use autoscaling if you wish to have a computational resource buffer in case of unexpected data surges. For this choose the minimum number of workers based on the projected daily data load and requirements and simply keep a couple of additional workers (say, 2 to 4) as a buffer to be added by autoscaling if and when needed.
- Workflows in Databricks allow you to share a job cluster with many tasks (jobs) that are all part of the same pipeline. If many jobs are executing in parallel on a shared job cluster, autoscaling for that job cluster should be enabled to allow it to scale up and supply resources to all of the parallel jobs.
- Autoscaling should not be used for Spark Structured Streaming workloads because once the cluster scales up, it is tough to determine when to scale it down again. There were some advances made in this aspect in Delta Live Tables, and autoscaling works pretty well for streaming workloads in Delta Live Tables than the feature called [Enhanced Autoscaling](#).

Instance types based on workloads

General guidelines to choose instance family based on the workload type:

VM Category	Workload Type
Memory Optimized	<ul style="list-style-type: none"> • For ML workloads • Where a lot of shuffle and spills are involved • When Spark caching is needed
Compute Optimized	<ul style="list-style-type: none"> • Structured Streaming jobs • ELT with full scan and no data reuse • To run OPTIMIZE and Z-order Delta commands
Storage Optimized	<ul style="list-style-type: none"> • To leverage Delta caching • ML and DL workloads with data caching • For ad hoc and interactive data analysis
GPU Optimized	<ul style="list-style-type: none"> • ML and DL workloads with an exceptionally high memory requirement
General Purpose	<ul style="list-style-type: none"> • Used in absence of specific requirements • To run VACUUM Delta command

Number of workers

Choosing the right number of workers requires some trials and iterations to figure out the compute and memory needs of a Spark job. Here are some guidelines to help you start:

- Never choose a single worker for a production job, as it will be the single point for failure
- Start with 2-4 workers for small workloads (for example, a job with no wide transformations like joins and aggregations)
- Start with 8-10 workers for medium to big workloads that involve wide transformations like joins and aggregations, then scale up if necessary
- Fine-tune the shuffle partitions when applicable to fully engage all cluster cores
- The approximate execution time can be determined after a few first executions. If this violates the SLA, you should add more workers.
- Remember that if shuffle partitions are fine-tuned and data skew and spill issues are addressed, adding additional workers will not result in a higher cost situation, adding more workers will proportionally reduce the total execution time, resulting in a cost that is more or less the same.

Spot instances

- Use spot instances to use spare VM instances for a below-market rate
- Great for interactive ad hoc/shared clusters
- Can use them for production jobs without SLAs
- Not recommended for production jobs with tight SLAs

- Never use spot instances for the driver

Auto-termination

- Terminates a cluster after a specified inactivity period in minutes for cost savings
- Enable [auto-termination](#) for all-purpose clusters to prevent idle clusters from running overnight or on weekends
- Clusters are configured to auto-terminate in 120 minutes by default, but you can change this to a much lower value, such as 10-15 minutes, to further save costs

Cluster usage

We need to make sure we're getting the most out of our cluster. A job may run many iterations on a cluster to complete depending on the number of shuffle partitions, so the rule of thumb is to ensure that all worker cores are actively occupied and not idle in any of the iterations. The only way to be absolutely certain is to always set the number of shuffle partitions as a multiplication of the total number of worker cores.

PYTHON (AUTO-DETECTED)

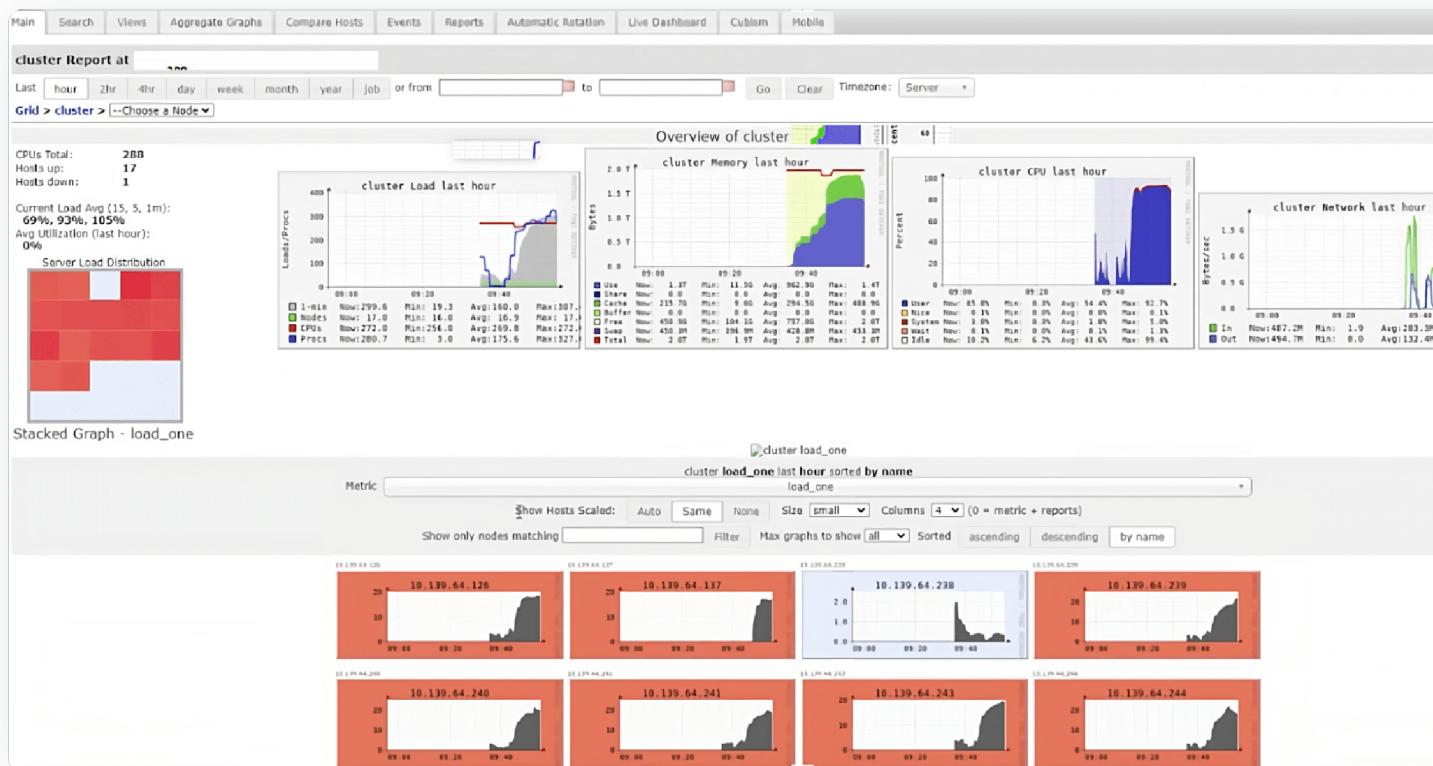
```
-- in SQL
set spark.sql.shuffle.partitions = M*<number of total cores in cluster>

-- in PySpark
spark.conf.set("spark.sql.shuffle.partitions", M*sc.defaultParallelism)

-- M is a multiplier here, pls refer manual shuffle partition tuning section above for more details

-- In the absence of manual shuffle partition tuning set M to 2 or 3 as a rule of thumb
```

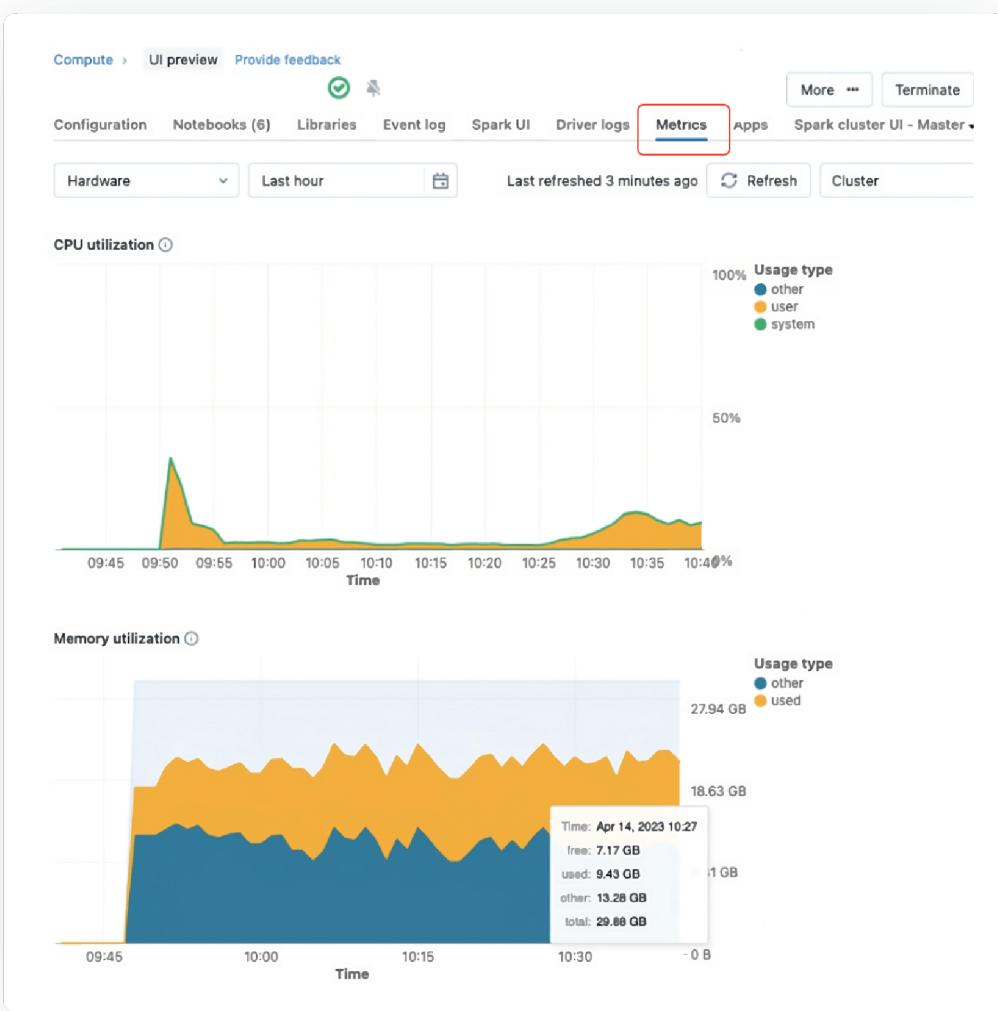
Ganglia UI is your friend to make sure that all the cores are fully engaged. Ganglia UI can be accessed right from the cluster UI pinned under the metrics tab



Things to pay attention to in the Ganglia UI:

- The average cluster load should always be greater than 80%
- During query execution, all squares in the cluster load distribution graph (on the left in the UI) should be red (with the exception of the driver node), indicating that all worker cores are fully engaged
- The use of cluster memory should be maximized (at least 60%-70%, or even more)
- Ganglia metrics are only available for Databricks Runtime 12.2 and below

Ganglia has been replaced in Databricks Runtime 13 and above by a new cluster metrics UI that is more secure and powerful, and provides a simple user experience with a clean design. Therefore, in DBR 13 and above, you can leverage this new UI to check out the cluster utilization.



Cluster policy

A [cluster policy](#) limits the ability to [configure clusters](#) based on a set of rules. Effective use of cluster policies allows administrators to:

- Limit users to create clusters with prescribed settings
- Control cost by limiting per-cluster maximum cost (by setting limits on attributes whose values contribute to hourly price)
- Simplify the user interface and enable more users to create their own clusters (by fixing and hiding some values)

Refer to the [cluster policies best practice guide](#) to plan and ensure a successful cluster governance rollout.

Instance pools

[Databricks pools](#) reduce cluster start and autoscaling times by maintaining a set of idle, ready-to-use instances. When a cluster is attached to a pool, cluster are created using the pool's idle instances. When instances are idle in the pool, they only incur Azure VM costs and no DBU costs.

Pools are recommended for workloads with tight SLAs, where fast access to additional compute resources is a requirement (i.e., fast autoscaling) to improve processing time while minimizing cost.

Photon

[Photon](#) is the next-generation engine on the Databricks Lakehouse Platform that provides extremely fast query performance at a low cost. Photon is compatible with Apache Spark APIs, works out of the box with existing Spark code, and provides significant performance benefits over the standard Databricks Runtime. Following are some of the advantages of Photon:

- Accelerated queries that process a significant amount of data (> 100GB) and include aggregations and joins
- Faster performance when data is accessed repeatedly from the Delta cache
- More robust scan/read performance on tables with many columns and many small files
- Faster Delta writing using UPDATE, DELETE, MERGE INTO, INSERT, and CREATE TABLE AS SELECT
- Join improvements

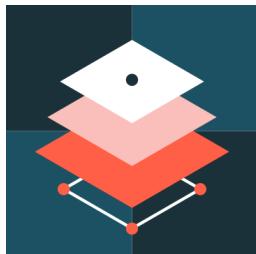
It is recommended to enable Photon (after evaluation and performance testing) for workloads with the following characteristics:

- ETL pipelines consisting of Delta MERGE operations

- Writing large volumes of data to cloud storage (Delta/Parquet)
- Scans of large data sets, joins, aggregations and decimal computations
- Auto Loader to incrementally and efficiently process new data arriving in storage
- Interactive/ad hoc queries using SQL

Others

- Use the latest LTS version of Databricks Runtime (DBR), as the latest Databricks Runtime is almost always faster than the one before it
- Restart all-purpose clusters periodically, at least once per week (or even daily on busy clusters), as some older DBRs might have stuck/zombie Spark jobs
- Configure a large driver node (4-8 cores and 16-32GB is mostly enough) only if:
 - Large data sets are being returned/collected (`spark.driver.maxResultSize` is typically increased also)
 - Large broadcast joins are being performed
- Many (50+) streams or concurrent jobs/notebooks on the same cluster:
 - Delta Live Tables might be a better fit for such workloads, as DLT determines the complete DAG of the pipeline and then goes about running it most optimized way possible
- For workloads using single-node libraries (e.g., Pandas), it is recommended to use **single-node data science clusters** as such libraries do not utilize the resources of a cluster in a distributed manner
- **Cluster tags** can be associated with Databricks clusters to attribute costs to specific teams/departments as they are automatically propagated to underlying cloud resources (VMs, storage and network, etc.)
- For such shuffle-heavy workloads, it is recommended to use fewer larger node sizes, which will help with reducing network I/O when transferring data between nodes
- Choosing workers with a large amount of RAM (>128GB) can help jobs perform more efficiently but can also lead to long pauses during **garbage collection**, which negatively impacts the performance and in some cases can cause job failure. Therefore, it's not recommended to choose workers with more than 128GB of RAM. Apart from worker memory, there are some other ways to avoid long garbage collection (GC) pauses:
 - Don't call `collect()` functions to collect data on the driver
 - Applies to the `toPandas()` function as well
 - Prefer Delta cache over Spark cache
 - If none of the above solutions help, use **G1GC** garbage collector instead by setting `spark.executor.defaultJavaOptions` and `spark.executor.extraJavaOptions` to `-XX:+UseG1GC` value in Spark config section under / cluster options
- Please refer to the [Cluster sizing examples](#) in Databricks documentation for specific cluster sizing examples



Conclusion

We have covered many Spark, Delta and Databricks optimization strategies and tactics in this book. The purpose of this book is to make readers aware of difficult challenges that they could encounter when developing big data workloads to run on distributed compute, as well as the methods for identifying these difficult solutions for them. A guide like this may truly assist data engineers and data architects in taking control of the cost of engineering and SLAs and making the use of cloud resources in the present economic climate, where every enterprise is attempting to cut costs and do more with less.